

Intelligent IT Support Ticket Classification and Response System using RAG

Project Proposal

1. Executive Summary

This project aims to design and deploy an intelligent IT support automation system that classifies incoming support tickets and generates accurate, context-aware responses using Retrieval-Augmented Generation (RAG).

The system integrates natural language processing (NLP), vector search, and transformer-based large language models (LLMs) to retrieve relevant historical IT incidents and produce grounded responses.

The final solution will be deployed as a scalable cloud API on Microsoft Azure and monitored using MLOps practices to ensure reliability, accuracy, and continuous improvement.

2. Project Objectives

The primary objectives of this project are:

- Automatically classify IT support tickets into predefined categories
 - Retrieve semantically similar historical IT issues from a knowledge base
 - Generate context-aware technical responses grounded in retrieved data
 - Deploy the system as a production-ready Azure API
 - Implement monitoring, logging, and retraining pipelines
-

3. Project Scope

In Scope

- IT support ticket dataset ingestion and preprocessing
- Text embeddings generation and vector database indexing
- Retrieval-Augmented Generation (RAG) pipeline implementation
- Transformer-based classification and response generation
- Azure cloud deployment (API + vector search integration)
- Monitoring dashboard and retraining triggers

Out of Scope

- Real company proprietary data integration
 - Multilingual support
 - Voice or chatbot interface
 - On-premise deployment
-

4. System Overview

The system processes IT support tickets through the following pipeline:

1. Ticket ingestion and preprocessing
2. Embedding generation
3. Vector similarity retrieval
4. Context injection into LLM
5. Response generation
6. API delivery to user

This architecture enables responses grounded in real IT incidents rather than generic model knowledge.

Project 3 (3)

5. Target Users

- IT helpdesk teams
 - Technical support agents
 - Enterprise IT service desks
 - Managed service providers (MSPs)
-

6. Expected Deliverables

By July 2026, the project will deliver:

- Cleaned and structured IT support ticket dataset
 - Vector search index and embeddings repository
 - Trained classification and RAG models
 - Deployed Azure API endpoint
 - Monitoring dashboard with KPIs
 - Final technical documentation and presentation
-

7. Technology Stack

NLP & ML

- Python
- HuggingFace Transformers
- Sentence-Transformers
- FAISS / Azure Cognitive Search

Backend & API

- FastAPI
- REST APIs

Cloud & MLOps

- Microsoft Azure ML

- Azure Cognitive Search
 - MLflow
 - Monitoring Dashboard
-

8. Project Timeline

Project duration: January 2026 → July 2026 (6 months)

Milestones:

1. Data Collection & Preprocessing
 2. Model Development (RAG)
 3. Azure Deployment
 4. MLOps & Monitoring
 5. Final Documentation & Demo
-

9. Project Team

Team size: 6 members

- محمد ابراهيم سعد
 - ايمان موسى محمود
 - هنا محمد وصفى
 - محمد أشرف محمد
 - احمد عبدالله ابراهيم
 - أشرف معوض رمضان
-

10. Success Criteria

The project will be considered successful if:

- Ticket classification accuracy meets target threshold
- Retrieved contexts are relevant to user queries

- Generated responses are technically correct and grounded
- API latency meets real-time requirements
- System runs reliably in Azure cloud environment

Project Plan

1. Official Project Timeline (DEPI)

Phase	Period	Description
Project Planning & Management	Jan → 20 Feb 2026	Proposal, plan, roles, risks, KPIs
Literature & Requirements	21 Feb → 20 Apr 2026	Research + requirements
System Analysis & Design	21 Apr → 1 May 2026	Architecture & diagrams
Implementation	2 May → 10 Jul 2026	RAG system development
Testing & Final Delivery	11 Jul → 17 Jul 2026	Testing + presentation

2. Milestones & Deliverables (Aligned with DEPI)

Milestone 1 — Project Planning & Management

Deadline: 20 Feb 2026

Deliverables:

- Project Proposal
- Project Plan
- Roles & Tasks
- Risk Plan
- KPIs

Milestone 2 — Literature Review & Requirements

Deadline: 20 Apr 2026

Deliverables:

- RAG & NLP literature review
 - Stakeholder analysis
 - User stories
 - Functional requirements
 - Non-functional requirements
-

Milestone 3 — System Analysis & Design

Deadline: 1 May 2026

Deliverables:

- System architecture diagram
 - Use case diagram
 - Data flow diagram
 - Component diagram
 - Deployment diagram
-

Milestone 4 — Implementation

Deadline: 10 Jul 2026

Deliverables:

- Data preprocessing pipeline
 - Embedding & vector DB
 - RAG retrieval system
 - LLM generation
 - Azure deployment
 - Source code (GitHub)
-

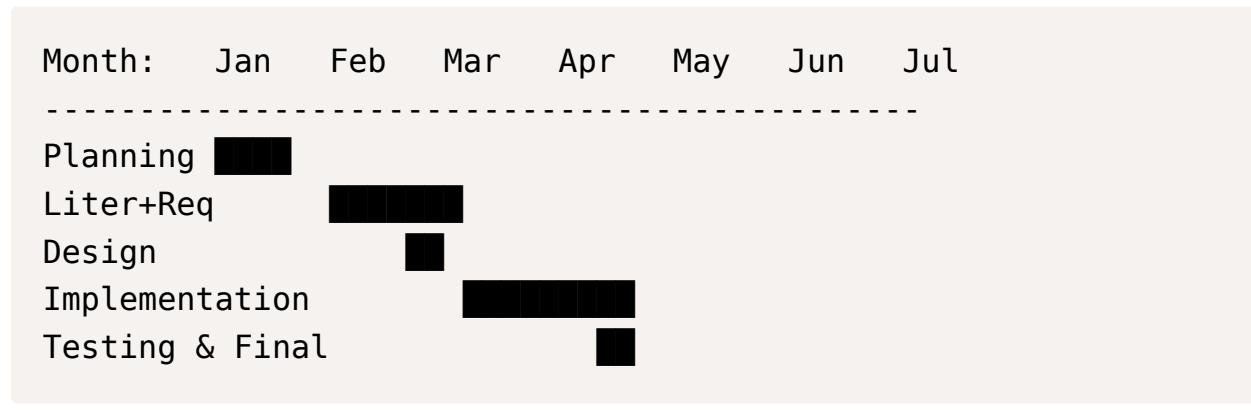
Milestone 5 — Testing & Final Delivery

Deadline: 17 Jul 2026

Deliverables:

- Test cases & results
- Performance evaluation
- Final report
- Presentation slides
- System demo

3. Student Gantt Chart (DEPI Schedule)



4. Work Distribution (Even Across Students)

Since DEPI requires each member contribution:

- Every phase → all 6 members contribute
- Tasks divided equally
- Deliverables merged collaboratively

Example distribution per phase

Phase	Work Split
Planning	1 doc per member
Literature	papers divided
Requirements	features divided
Design	diagrams divided

Phase	Work Split
Implementation	modules divided
Testing	cases divided

5. Phase Dependencies

DEPI workflow is sequential:

Planning → Requirements → Design → Implementation → Testing

Implementation cannot start before design approval.

6. Key Submission Dates (DEPI)

Deliverable	Date
Planning Docs	20 Feb 2026
Literature & Requirements	20 Apr 2026
Design	1 May 2026
Implementation	10 Jul 2026
Final Presentation	17 Jul 2026

Task Assignment & Roles

1. Team Structure (XP Pairs)

Pair	Members
Pair 1	إيمان موسى محمود — هنا محمد وصفى
Pair 2	محمد ابراهيم سعد — احمد عبدالله ابراهيم
Pair 3	محمد أشرف محمد — أشرف معوض رمضان

Each pair works together on all tasks through continuous collaboration, code review, and shared ownership, following XP principles.

2. XP Role Model in Project

In Extreme Programming, roles rotate rather than being fixed.

Each pair alternates between:

- Driver (coding)
- Navigator (reviewing & guiding)

This ensures equal contribution and knowledge sharing.

3. Responsibilities by Pair (Project Phases)

Pair 1 — Data & NLP Pipeline

Members: ايمان موسى محمود — هنا محمد وصفى

Responsibilities:

- Dataset collection & cleaning
- Text preprocessing
- Tokenization & normalization
- Embedding generation
- EDA analysis

Deliverables:

- Clean dataset
 - Preprocessing scripts
 - EDA report
-

Pair 2 — RAG Retrieval & Modeling

Members: محمد ابراهيم سعد — احمد عبدالله ابراهيم

Responsibilities:

- Vector database construction
- FAISS / Azure Cognitive Search indexing

- Similarity search implementation
- Ticket classification model
- Retrieval evaluation

Deliverables:

- Vector index
 - Retrieval module
 - Classification model
-

Pair 3 — Generation, API & Deployment

Members: محمد أشرف محمد — أشرف معاوض رمضان

Responsibilities:

- LLM prompt design
- RAG response generation
- FastAPI backend
- Azure deployment
- System integration

Deliverables:

- RAG generation module
 - REST API
 - Azure deployment
-

4. Shared Responsibilities (All Pairs)

All pairs collaborate on:

- System architecture design
- Integration testing
- Evaluation & metrics

- Documentation
- Final presentation

This satisfies DEPI requirement that each student contributes to the full project lifecycle.

Project Documentation (1)

5. XP Practices Applied in Project

The team follows Extreme Programming practices:

- Pair programming
 - Continuous integration
 - Small iterative releases
 - Refactoring
 - Collective code ownership
 - Frequent testing
-

6. Workload Balance Assurance

Balance is ensured by:

- Equal pair size (2 each)
- Rotating driver/navigator roles
- Cross-phase collaboration
- Shared final deliverables

Thus, all six members contribute equally to planning, development, and delivery.

Risk Assessment & Mitigation Plan

1. Risk Management Approach

The team follows proactive risk management:

- Early identification
- Probability & impact evaluation
- Preventive mitigation
- Monitoring during milestones

Risks are reviewed at the end of each project phase.

2. Risk Matrix

ID	Risk	Probability	Impact	Level
R1	Poor data quality	Medium	High	High
R2	Weak embedding/retrieval accuracy	Medium	High	High
R3	Azure deployment complexity	Medium	Medium	Medium
R4	LLM hallucination	Medium	High	High
R5	Integration failures	Low	High	Medium
R6	Uneven contribution in pairs	Low	Medium	Low
R7	Time constraints before deadlines	Medium	High	High
R8	System latency (slow response)	Medium	Medium	Medium

3. Risk Details & Mitigation

R1 — Poor Data Quality

Description: IT tickets may contain noise, missing labels, or inconsistent text.

Impact: Weak training and retrieval performance.

Mitigation:

- Text cleaning pipeline

- Manual inspection samples
- Remove duplicates
- Normalize categories

Owner: Pair 1 (Data)

R2 — Weak Retrieval Accuracy

Description: Retrieved tickets may not match user query semantics.

Impact: Incorrect generated responses.

Mitigation:

- Use sentence-transformer embeddings
- Tune similarity metric & top-k
- Evaluate retrieval accuracy
- Compare models

Owner: Pair 2 (Retrieval)

R3 — Azure Deployment Complexity

Description: Cloud services integration may fail or be misconfigured.

Impact: System not deployable.

Mitigation:

- Local prototype first
- Use Azure tutorials & templates
- Incremental deployment
- Versioned configuration

Owner: Pair 3 (Deployment)

R4 — LLM Hallucination

Description: Model generates unsupported or incorrect IT advice.

Impact: Loss of reliability.

Mitigation:

- Strict RAG context grounding
- Limit generation temperature
- Include retrieved evidence
- Human validation samples

Owner: Pair 3 (Generation)

R5 — Integration Failures

Description: Retrieval, LLM, and API modules may not connect correctly.

Impact: End-to-end pipeline breaks.

Mitigation:

- Modular interfaces
- Integration testing early
- API contracts defined
- Continuous integration

Owner: All pairs

R6 — Uneven Contribution in XP Pairs

Description: One student may dominate coding tasks.

Impact: Unfair workload & grading issues.

Mitigation:

- Driver/navigator rotation
- Weekly role switching
- GitHub commit tracking
- Pair reviews

Owner: All pairs

R7 — Time Constraints

Description: Implementation may exceed DEPI deadlines.

Impact: Late submission.

Mitigation:

- Phase-based milestones
- MVP first approach
- Weekly progress checks
- Scope control

Owner: All pairs

R8 — High Response Latency

Description: RAG pipeline may respond slowly due to vector search or LLM.

Impact: Poor user experience.

Mitigation:

- Limit context size
- Optimize top-k
- Cache embeddings
- Measure response time KPI

Owner: Pair 2 & 3

4. Risk Monitoring Plan

Risks reviewed at:

- Phase completion meetings
- Before each DEPI submission
- Integration milestones

If risk level increases → mitigation updated.

5. Contingency Strategy

If major technical risk occurs:

- Switch to smaller embedding model
- Replace Azure with FAISS local
- Use simpler classifier baseline
- Reduce dataset size

Ensures project completion within deadline.

Key Performance Indicators (KPIs)

1. KPI Categories

Project success is evaluated across five dimensions:

1. Retrieval Quality
2. Classification Accuracy
3. Response Generation Quality
4. System Performance
5. Project Delivery & Usability

2. Retrieval KPIs (RAG Core)

KPI 1 — Top-K Retrieval Accuracy

Definition: Percentage of queries where relevant ticket appears in top-K results

Formula:

Relevant_in_topK / Total_queries

Target: $\geq 80\%$

Why: Measures vector search effectiveness.

KPI 2 — Similarity Score Quality

Definition: Average cosine similarity between query and retrieved tickets

Target: ≥ 0.70

Why: Ensures semantic closeness of retrieved context.

KPI 3 — Retrieval Latency

Definition: Time to retrieve top-K documents

Target: ≤ 300 ms

Why: Important for real-time API.

3. Classification KPIs

KPI 4 — Ticket Classification Accuracy

Definition: Correct category predictions / total tickets

Target: $\geq 85\%$

KPI 5 — Recall (Critical Classes)

Definition: Ability to detect key IT issue categories

Target: $\geq 80\%$

Why: Missing important IT issues is costly.

4. Generation KPIs (LLM + RAG)

KPI 6 — Response Groundedness

Definition: % of responses supported by retrieved tickets

Measurement: Manual evaluation sample

Target: $\geq 85\%$

Why: Prevent hallucination.

KPI 7 — Response Relevance Score

Definition: Human rating (1–5) of answer usefulness

Target: ≥ 4.0

KPI 8 — BLEU / ROUGE (Optional)

Definition: Overlap with known resolutions

Target: Moderate (reference only)

Note: Human evaluation preferred for RAG text.

5. System Performance KPIs

KPI 9 — End-to-End Response Time

Definition: Time from ticket input → final response

Target: ≤ 2 seconds

KPI 10 — API Availability

Definition: % uptime of Azure endpoint

Target: $\geq 95\%$

KPI 11 — System Stability

Definition: Failure rate during requests

Target: $\leq 5\%$ errors

6. Project Execution KPIs (DEPI)

KPI 12 — Milestone Completion Rate

Definition: Deliverables completed on deadline

Target: 100%

KPI 13 — Code Contribution Balance

Definition: Even GitHub commits across members

Target: Balanced across 6 students

KPI 14 — Documentation Completeness

Definition: All DEPI required documents submitted

Target: 100%

7. KPI Summary Table

Category	KPI	Target
Retrieval	Top-K accuracy	$\geq 80\%$
Retrieval	Similarity	≥ 0.70
Classification	Accuracy	$\geq 85\%$
Classification	Recall	$\geq 80\%$
Generation	Groundedness	$\geq 85\%$
Generation	Relevance	$\geq 4/5$
Performance	Response time	$\leq 2 \text{ s}$
Performance	API uptime	$\geq 95\%$
Project	Milestones	100%
Project	Contribution	Balanced

8. KPI Measurement Plan

KPIs measured at:

- After model training
- After deployment
- Before final submission

Tools:

- sklearn metrics
- retrieval evaluation scripts

- latency logging
- human evaluation forms
- GitHub analytics